

SEYOUNG OH

San Jose, CA | 408-304-8254 | syoh2k@gmail.com | [linkedin.com/in/syoh2k](https://www.linkedin.com/in/syoh2k) | github.com/driffe

SUMMARY

Full-stack engineer specializing in production LLM systems. Built and ran a live AI platform end-to-end with a RAG engine, multi-agent orchestration with human review gates, and offline evals. Backend depth from a 25x throughput rewrite and shipping to over 900K monthly users at KQED. 2x hackathon winner.

EXPERIENCE

KQED

San Francisco, CA

Junior Software Engineer

Mar 2026 - Present

- Debugged and resolved a critical production bug on KQED.org (900K+ monthly visitors), tracing the root cause across server-side rendering and frontend caching layers to restore full page functionality.
- Built end-to-end pagination for 9,800+ tag-archive pages on KQED.org, implementing CSR/SSR canonical redirects, dynamic SEO tags from Redux state, VoiceOver accessibility, and 404 handling for invalid requests.

Software Engineer Intern

Jul 2025 - Feb 2026

- Increased a file upload service's capacity by 25x (from 2 to 50+ requests/sec) by rewriting blocking database calls into a non-blocking request pipeline, eliminating timeouts that previously occurred under heavy load.
- Cut the team's API maintenance load by removing 27% of unused endpoints after auditing 100+ APIs across services in Postman within a microservices architecture (Spring WebFlux) using Kafka for event streaming.

JKI INC

Campbell, CA

Field Application Engineer

Sep 2024 - Jan 2025

- Embedded at a customer site as the primary technical liaison between the customer's engineering team and JKI's product teams, supporting hardware qualification of NAND flash memory products.
- Built a C++ test automation tool that runs full hardware validation with one command, and discovered an additional 17% safe operating voltage range that informed the customer's production qualification spec.

PROJECTS

Coffeeyme | Spring Boot, Next.js, Python, AWS (Bedrock, EC2, RDS), GitHub Actions, Docker

Technical Lead

Jul 2025 - Present

- Built a RAG recommendation coffee platform over 700+ products using AWS Bedrock and Amazon Titan embeddings, running a hard filter, TopN retrieval, Claude final selection, with Redis caching for sub-100ms repeat queries.
- Built a multi-agent marketing system where a manager agent plans work, gates for human review, and dispatches to 5 specialized agents covering keyword research, SEO, SNS content, and roaster profiles. Added a Python verification layer that automatically fails the pipeline on data leakage between agents.
- Built a cost-optimized LLM data pipeline scraping 40+ roaster catalogs via Shopify API with Selenium fallback, using a 2-tier Batch API flow with SHA-256 idempotency to cut AI cost, and validated output quality with an offline RAG eval pipeline for regression testing.

AWARDS

1st Place - AI Sports, IBM SkillsBuild: NBA AI Platform with LLM chatbot (FastAPI + IBM Watsonx Granite) (2025)

1st Place - Perplexity, HackHayward: AI Medication Recommendation System (FastAPI + Perplexity) (2025)

TECHNICAL SKILLS

Languages: Java, Python, JavaScript, TypeScript, C++, SQL, Swift

Frameworks: Spring Boot, Spring WebFlux, FastAPI, Node.js, React, Next.js

AI / LLM: AWS Bedrock, Anthropic Claude API, OpenAI API, RAG Pipelines, Semantic Search, Vector Embeddings, Prompt Engineering, Eval Design, Multi-Agent Orchestration

Cloud & DevOps: AWS (EC2, RDS, Bedrock, S3), GCP, Docker, Nginx, GitHub Actions

Databases & Tools: MySQL, MongoDB, Firebase, Redis, REST APIs, Git, Postman, Jira

EDUCATION

San Jose State University

B.S. in Software Engineering

San Jose, CA

Aug 2021 - Dec 2025